A study of ACO capabilities for solving the Maximum Clique Problem

Article *in* Journal of Heuristics · January 2006

2 authors, including:

Christine Solnon
Institut National des Sciences Appliquées de Lyon
90 PUBLICATIONS **1,741** CITATIONS

Some of the authors of this publication are also working on these related projects:

Car sequencing View project

A study of ACO capabilities for solving the Maximum Clique Problem

Christine Solnon and Serge Fenet
LIRIS CNRS UMR 5205, University Lyon I
Nautibus, 43 Bd du 11 novembre, 69622 Villeurbanne cedex, France
{christine.solnon,serge.fenet}@liris.cnrs.fr

June 23, 2005

Abstract

This paper investigates the capabilities of the Ant Colony Optimization (ACO) meta-heuristic for solving the maximum clique problem, the goal of which is to find a largest set of pairwise adjacent vertices in a graph. We propose and compare two different instantiations of a generic ACO algorithm for this problem. Basically, the generic ACO algorithm successively generates maximal cliques through the repeated addition of vertices into partial cliques, and uses "pheromone trails" as a greedy heuristic to choose, at each step, the next vertex to enter the clique. The two instantiations differ in the way pheromone trails are laid and exploited, i.e., on edges or on vertices of the graph.

We illustrate the behavior of the two ACO instantiations on a representative benchmark instance and we study the impact of pheromone on the solution process. We consider two measures —the re-sampling and the dispersion ratio— for providing an insight into the performance at run time. We also study the benefit of integrating a local search procedure within the proposed ACO algorithm, and we show that this improves the solution process. Finally, we compare ACO performance with that of three other representative heuristic approaches, showing that the former obtains competitive results.

1 Introduction

The maximum clique problem is a classical combinatorial optimization problem that has important applications in different domains such as coding theory, fault diagnosis, or computer vision. Moreover, many important problems —such as constraint satisfaction, subgraph isomorphism, or vertex covering problems— are easily reducible to this maximum clique problem.

Given a non-directed graph $G = (V, E)$, such that V is a set of vertices, and $E \subseteq V \times V$ is a set of edges, a *clique* is a set of vertices $\mathcal{C} \subseteq V$ such that every

pair of distinct vertices of $\mathcal{C}$ is connected with an edge in G, i.e., the subgraph induced by $\mathcal{C}$ is complete. A clique is *partial* if it is strictly included in another clique; otherwise it is *maximal*. The goal of the maximum clique problem is to find a clique of maximum cardinality.

The maximum clique problem is very closely related to the maximum independent (or stable) set problem, the goal of which is to find the maximum subset of V such that no two vertices of the subset are pairwise adjacent: a maximum clique of a graph $G = (V, E)$ is a maximum independent set of the complement graph $G' = (V, V \times V - E)$ of G.

The clique problem —that consists in finding a clique of size k in a graph— is one of the first problems shown to be NP-complete in [1]. More generally, the problem of finding a maximum clique is NP-hard, and does not admit polynomial-time approximation algorithms (unless P=NP) [2]. Hence, exact approaches —usually based on branch-and-bound— become intractable when the number of vertices increases, and much effort has recently been directed on heuristic approaches.

Heuristic approaches for the maximum clique problem

Heuristic approaches leave out exhaustivity and use heuristics to guide the search towards promising areas of the search space. They run in polynomial time and quickly find "rather good" solutions, that may be optimal, but optimality is not guaranteed.

Many heuristic approaches are based on sequential greedy heuristics [3, 4, 5, 6, 7, 8]. The idea is to build maximal cliques, starting from an empty clique, and iterating through the repeated addition of vertices. To decide which vertex is added, one uses a greedy heuristic such as, e.g., choosing the vertex that has the highest degree among candidate vertices. To avoid usual greedy traps, greedy heuristics can be improved by injecting a mild amount of randomization combined with multiple restarts. Also, weights used by the greedy heuristic may be adapted from restart to restart as proposed, e.g., in [7, 8].

To improve the quality of a constructed clique, local search can be used to explore its neighborhood, i.e., the set of cliques that can be obtained by removing and/or adding a given number of vertices: local search iteratively moves in the search space composed of all cliques, from a clique to one of its (best) neighbors. To avoid being trapped in local optima, where all neighbors are cliques of smaller sizes, local search may be combined with some advanced meta-heuristics. For example in [9, 10], Simulated Annealing is used to jump out of local optima by allowing moves towards smaller cliques with a probability proportional to a decreasing temperature parameter. In [11, 12, 13], Tabu Search is used to prevent local search from cycling through a small set of good but suboptimal cliques by keeping track, in a tabu list, of forbidden moves between cliques. In [6], Reactive Search enhances Tabu Search by reactively adapting the size of the tabu list with respect to the need of diversification. In [14], local search is combined with a Genetic Algorithm that allows one to escape from local maxima by applying crossover and mutation operators to a population of

maximal cliques.

The Ant Colony Optimization meta-heuristic

In this paper, we investigate the capabilities of another meta-heuristic —Ant Colony Optimization (ACO) [15, 16]— for solving the maximum clique problem. The basic idea of ACO is to model the problem to solve as the search for a minimum cost path in a graph, and to use artificial ants to search for best paths.

The behavior of artificial ants is inspired from real ants: artificial ants lay pheromone trails on components of the graph and choose their path with respect to probabilities that depend on pheromone trails that have been previously laid; these pheromone trails progressively decrease by evaporation. Intuitively, this indirect stigmergic communication means aims at giving information about the quality of path components in order to attract ants, in the following iterations, towards the corresponding areas of the search space. Indeed, for many combinatorial problems, a study of the search space landscape shows a correlation between solution quality and the distance to optimal solutions [17, 18, 19].

The first ant algorithm to be applied to a discrete optimization problem has been proposed by Dorigo in [20]. The problem chosen for the first experiments was the Traveling Salesman Problem and, since then, this problem has been widely used to investigate the solving capabilities of ants [21, 22]. The ACO meta-heuristic, described in [15, 16], is a generalization of these first ant based algorithms, and it has been successfully applied to different hard combinatorial optimization problems such as quadratic assignment problems [23, 24], vehicle routing problems [25, 26], and constraint satisfaction problems [27, 28].

We have proposed in [29] a first ACO algorithm for the maximum clique problem. The contribution of this paper with respect to this preliminary work mainly concerns (i) the definition of a variant of this first ACO algorithm, that differs in the way pheromone is laid and exploited; (ii) an experimental comparison of these two different ways of managing pheromone; (iii) the use of diversity measures in order to provide an insight into the performance at runtime; (iv) an investigation of the benefit of combining ACO with local search for this problem; and (v) an experimental comparison with three state-of-the-art heuristic approaches.

Another Ant algorithm for the maximum clique problem has been recently proposed in [30]. This algorithm is rather different as ants are distributed, i.e., each ant only has local knowledge of the graph, so that the algorithm can be implemented in a distributed system. As a counterpart, performance of this algorithm by means of solutions' quality are rather far from performance of the ACO algorithm proposed in [29] and other non distributed algorithms.

Overview of the paper

Section 2 describes the generic ACO algorithm for the maximum clique problem —called `Ant-Clique`— and its two different instantiations —called `Vertex-AC`

and `Edge-AC`. Basically, `Ant-Clique` uses pheromone trails as a greedy heuristic for choosing, at each step, the next vertex to enter the clique. However, in `Vertex-AC`, pheromone trails are laid on vertices, and the choice of vertices directly depends on the quantity of pheromone laying on them, whereas in `Edge-AC`, pheromone trails are laid on edges, and the choice of vertices depends on pheromone trails laying on edges connecting candidate vertices with the vertices of the clique under construction.

In Section 3, we illustrate and compare the behavior of `Vertex-AC` and `Edge-AC` on a representative benchmark instance. We consider two measures in order to provide a deeper insight into their behavior at run-time: the re-sampling ratio is used to measure how often the search space is re-sampled, whereas the dispersion ratio is used to quantify the differences between the cliques successively computed during the solution process.

In Section 4, we show how local search can be combined with `Ant-Clique`, and we study on a representative benchmark instance the influence of local search on the solution process.

Section 5 experimentally compares the four different ACO instantiations (`Vertex-AC` and `Edge-AC`, with and without local search) on a set of benchmark graphs. We show that `Edge-AC` usually finds better solutions than `Vertex-AC`, but is more time consuming. We also show that the integration of local search improves solutions' quality on a majority of instances.

Finally, Section 6 compares our approach with other heuristic approaches. We have more particularly chosen for this comparison three recent and representative algorithms within different classes of heuristic approaches: `RLS` [6], which uses reactive local search and obtains the best known results on most benchmark instances, `DAGS` [8], which combines an adaptive greedy approach with swap local moves and obtains the best known results on some other instances, and `GLS` [14], which combines a genetic approach with local search.

2 ACO for the maximum clique problem

In ACO algorithms, ants lay pheromone trails on components of the best constructed solutions in order to attract other ants towards the corresponding area of the search space. To solve a new problem with ACO, one mainly has to define the pheromone laying procedure —i.e., decide on which components of constructed solutions ants should lay pheromone trails— and define the solution construction procedure —i.e., decide how to exploit these pheromone trails when constructing new solutions.

Hence, to solve the maximum clique problem with ACO, the key point is to decide which components of the constructed cliques should be rewarded, and how to exploit these rewards when constructing new cliques. Indeed, given a maximal clique $\mathcal{C}_i$, one can lay pheromone trails either on the vertices of $\mathcal{C}_i$, or on the edges connecting every pair of different vertices of $\mathcal{C}_i$:

- when laying pheromone on the vertices of $\mathcal{C}_i$, the idea is to increase the

Search of an approximate maximum clique in a graph $G = (V, E)$:
    Initialize pheromone trails to τ_{max}
    repeat the following cycle:
        for each ant k **in** $1..nbAnts$, construct a maximal clique $\mathcal{C}_k$ as follows:
            Randomly choose a first vertex $v_i \in V$
            $\mathcal{C}_k \leftarrow \{v_i\}$
            $Candidates \leftarrow \{v_j \in V \mid (v_i, v_j) \in E\}$
            while $Candidates \neq \emptyset$ **do**
                Choose a vertex $v_i \in Candidates$ with probability $p(v_i)$
                $\mathcal{C}_k \leftarrow \mathcal{C}_k \cup \{v_i\}$
                $Candidates \leftarrow Candidates \cap \{v_j \mid (v_i, v_j) \in E\}$
            end while
        end for
        Update pheromone trails w.r.t. $\{\mathcal{C}_1, \ldots, \mathcal{C}_{nbAnts}\}$
        if a pheromone trail is lower than τ_{min} **then** set it to τ_{min}
        if a pheromone trail is greater than τ_{max} **then** set it to τ_{max}
    until maximum number of cycles reached **or** optimal solution found
    return the largest constructed clique since the beginning

Figure 1: Generic algorithmic scheme of `Ant-Clique`

desirability of each vertex of $\mathcal{C}_i$ so that, when constructing a new clique,
these vertices will be more likely to be selected;

- when laying pheromone on the edges of $\mathcal{C}_i$, the idea is to increase the desir-
  ability of choosing together two vertices of $\mathcal{C}_i$ so that, when constructing
  a new clique $\mathcal{C}_k$, the vertices of $\mathcal{C}_i$ will be more likely to be selected *if
  $\mathcal{C}_k$ already contains some vertices of $\mathcal{C}_i$*. More precisely, the more $\mathcal{C}_k$ will
  contain vertices of $\mathcal{C}_i$, the more the other vertices of $\mathcal{C}_i$ will be attractive.

A goal of this paper is to compare these two different ways of using pheromone.
Therefore, we introduce a generic ACO algorithm called `Ant-Clique`, and two
different instantiations of it: `Vertex-AC`, where pheromone is laid on vertices,
and `Edge-AC`, where pheromone is laid on edges.

In this section, we first describe the generic algorithmic scheme that is com-
mon to the two instantiations. Then, we describe the three points on which
they differ, i.e., the definition of pheromonal components, the exploitation of
pheromone trails when constructing cliques, and the pheromone updating pro-
cess. Finally, we compare their time complexities.

2.1   Generic algorithmic scheme of `Ant-Clique`

Figure 1 displays the generic ACO algorithmic scheme for solving the maxi-
mum clique problem. At each cycle of this algorithm, every ant constructs a
maximal clique. It first randomly chooses an initial vertex to enter the clique,

and then iteratively adds vertices that are chosen within a set *Candidates* that contains all the vertices that are connected to every vertex of the partial clique under construction. This choice is performed randomly with respect to probabilities that are defined in Section 2.3. Once each ant has constructed a clique, pheromone trails are updated, as described in Section 2.4. The algorithm stops either when an ant has found a maximum clique (when the optimal bound is known), or when a maximum number of cycles has been performed.

This algorithm more particularly borrows features from the $\mathcal{MAX} - \mathcal{MIN}$ Ant System [19]: it explicitly imposes lower and upper bounds τ_{min} and τ_{max} on pheromone trails (with $0 < \tau_{min} < \tau_{max}$), and pheromone trails are set to τ_{max} at the beginning of the search.

2.2 Definition of pheromonal components

Pheromone trails are laid by ants on components of the graph $G = (V, E)$ in which they are looking for a maximum clique.

- In `Vertex-AC`, ants lay pheromone trails on the vertices V of the graph. The quantity of pheromone on a vertex $v_i \in V$ is denoted τ_i. Intuitively, this quantity represents the learned desirability to select v_i when constructing a clique.

- In `Edge-AC`, ants lay pheromone trails on the edges E of the graph. The quantity of pheromone on an edge $(v_i, v_j) \in E$ is denoted τ_{ij}. Intuitively, this quantity represents the learned desirability to select v_i when constructing a clique that already contains v_j. Notice that since the graph is not directed, $\tau_{ij} = \tau_{ji}$.

2.3 Exploitation of pheromone trails

Pheromone trails are used to choose vertices when constructing cliques: at each step, a vertex v_i is randomly chosen within the set *Candidates* with respect to a probability $p(v_i)$. This probability is defined proportionally to pheromone factors, i.e.,

$$p(v_i) = \frac{[\tau fact(v_i)]^\alpha}{\sum_{v_j \in Candidates} [\tau fact(v_j)]^\alpha}$$

where α is a parameter which weights pheromone factors, and $\tau fact(v_i)$ is the pheromone factor of vertex v_i. This pheromone factor depends on the quantity of pheromone laying on pheromonal components of the graph:

- in `Vertex-AC`, it depends on the quantity of pheromone laid on the candidate vertex, i.e.,

$$\tau fact(v_i) = \tau_i$$

- in `Edge-AC`, it depends on the quantity of pheromone laid on edges connecting the vertices that already are in the partial clique and the candidate

6

vertex: let $\mathcal{C}_k$ be the partial clique under construction, the pheromone factor of a candidate vertex v_i is

$$\tau fact(v_i) = \sum_{v_j \in \mathcal{C}_k} \tau_{ij}$$

Note that this pheromone factor can be computed in an incremental way: once the first vertex v_i has been randomly chosen, for each candidate vertex v_j that is adjacent to v_i, the pheromone factor $\tau fact(v_j)$ is initialized to τ_{ij}; then, each time a new vertex v_k is added to the clique, for each candidate vertex v_j that is adjacent to v_k, the pheromone factor $\tau fact(v_j)$ is incremented by τ_{kj}.

One should remark that, for both `Vertex-AC` and `Edge-AC`, the probability of choosing a vertex v_i only depends on a pheromone factor and not on some other heuristic factor that locally evaluates the quality of the candidate vertex, as usually in ACO algorithms. Actually, we have experimented different heuristics, and more particularly the greedy heuristic proposed, e.g., in [3, 6, 8]. The idea is to favor vertices with largest degrees in the "residual graph", i.e., the subgraph induced by the set of candidate vertices that can extend the partial clique under construction. The underlying motivation is that a larger degree implies a larger number of candidates once the vertex has been added to the current clique. When using no pheromone, or at the beginning of the search process when all pheromone trails have the same value, this heuristic actually allows ants to construct larger cliques with respect to a random choice. However, when combining it with pheromone learning, we have noticed that after a hundred or so cycles we obtain larger cliques without using the heuristic than when using it.

To explain this rather counter-intuitive result, we have compared re-sampling ratio (i.e., the percentage of solutions that are re-computed several times [31]) at run-time. We have noticed that when integrating the greedy heuristic to the probability $p(v_i)$, this re-sampling ratio is significantly increased. This shows that the search is trapped around a set of locally optimal cliques, and is not diversified enough to be able to find larger cliques.

2.4 Updating pheromone trails

Once each ant has constructed a clique, the amount of pheromone laying on pheromonal components is updated according to the ACO meta-heuristic. First, all amounts are decreased in order to simulate evaporation. This is done by multiplying the quantity of pheromone laying on each pheromonal component by a pheromone persistence rate ρ such that $0 \le \rho \le 1$. Then, the best ant of the cycle deposits pheromone. More precisely, let $\mathcal{C}_k \in \{\mathcal{C}_1, \ldots, \mathcal{C}_{nbAnts}\}$ be the largest clique built during the cycle (if there are several largest cliques, ties are randomly broken), and $\mathcal{C}_{best}$ be the largest clique built since the beginning of the run (including the current cycle). The quantity of pheromone laid by ant k is inversely proportional to the gap of size between $\mathcal{C}_k$ and $\mathcal{C}_{best}$, i.e., it is equal

to $1/(1+ \mid \mathcal{C}_{\text{best}} \mid - \mid \mathcal{C}_k \mid)$. This quantity of pheromone is deposited on the pheromonal components of $\mathcal{C}_k$, i.e.,

- in `Vertex-AC`, it is deposited on each vertex of $\mathcal{C}_k$,

- in `Edge-AC`, it is deposited on each edge connecting two different vertices of $\mathcal{C}_k$.

2.5 Time complexity of `Vertex-AC` and `Edge-AC`

Let $nbMaxCycles$ and $nbAnts$ respectively be the maximum number of cycles and the number of ants. In the worst case (if an optimal solution is not found), both `Vertex-AC` and `Edge-AC` will have to construct $nbMaxCycles \cdot nbAnts$ maximal cliques, and to perform $nbMaxCycles$ pheromone updating steps.

To construct a maximal clique, `Vertex-AC` and `Edge-AC` nearly perform the same number of operations. Indeed, to construct a clique $\mathcal{C}$, both instantiations perform $\mid \mathcal{C} \mid -1$ times the "while" loop of the algorithm of Figure 1. At each iteration of this loop, they both have to: (i) compute probabilities for all candidates, (ii) choose a vertex with respect to these probabilities, and (iii) update the list of candidates. These three steps take a linear time with respect to the number of candidate vertices[1]. At the first iteration, the number of candidate vertices is equal to the degree of the initial vertex, and it decreases at each iteration. Hence, in the worst case, the complexity of the construction of a clique in a graph G, for both `Vertex-AC` and `Edge-AC`, is in $\mathcal{O}(\omega(G) \cdot max_{d^\circ}(G))$, where $\omega(G)$ is the size of the maximum clique of G, and $max_{d^\circ}(G)$ is the maximum vertex degree in G. Note that both $\omega(G)$ and $max_{d^\circ}(G)$ are bounded by the number of vertices of G.

To update pheromone trails laying on pheromonal components of a graph $G = (V, E)$, `Vertex-AC` and `Edge-AC` perform a different number of operations:

- In `Vertex-AC`, the evaporation step requires $\mathcal{O}(\mid V \mid)$ operations and the reward of a clique $\mathcal{C}$ requires $\mathcal{O}(\mid \mathcal{C} \mid)$ operations. As $\mid \mathcal{C} \mid \leq \mid V \mid$, the whole pheromone updating step requires $\mathcal{O}(\mid V \mid)$ operations.

- In `Edge-AC`, the evaporation step requires $\mathcal{O}(\mid E \mid)$ operations and the reward of a clique $\mathcal{C}$ requires adding pheromone on each edge of $\mathcal{C}$. As the number of edges of a clique is smaller or equal to $\mid E \mid$, the whole pheromone updating step requires $\mathcal{O}(\mid E \mid)$ operations.

Hence, the overall time complexity of `Vertex-AC` is in

$$\mathcal{O}(nbMaxCycles(nbAnts \cdot \omega(G) \cdot max_{d^\circ}(G) + \mid V \mid))$$

whereas the overall time complexity of `Edge-AC` is in

$$\mathcal{O}(nbMaxCycles(nbAnts \cdot \omega(G) \cdot max_{d^\circ}(G) + \mid E \mid))$$

---

[1]Remember that, as pointed out in 2.3, pheromone factors in `Edge-AC` are computed in an incremental way, i.e., each time a new vertex is added to the clique, the pheromone factor of each candidate vertex is updated by a simple addition.

3 Diversification versus intensification of search

When solving a combinatorial optimization problem with a heuristic approach such as evolutionary computation or ACO, one usually has to find a compromise between two dual goals. On the one hand, one has to intensify the search around the most "promising" areas, that are usually close to the best solutions found so far [17, 18, 19]. On the other hand, one has to diversify the search and favor exploration in order to discover new, and hopefully more successful, areas of the search space. The behavior of ants with respect to this intensification/diversification duality can be influenced by modifying parameter values [32].

In this section, we first briefly discuss the settings of τ_{min}, τ_{max}, and $nbAnts$. Then, we study the influence of α and ρ on the solution process on a representative benchmark instance. Finally, we provide an insight into the influence of α and ρ on the intensification/diversification duality by means of two diversity measures.

3.1 Bounding pheromone trails within $[\tau_{min}, \tau_{max}]$

As pointed out in [19], the goal of bounding pheromone trails within an interval $[\tau_{min}, \tau_{max}]$ is to avoid premature stagnation of search, i.e., a situation where all ants construct the same solution over and over again so that no better solutions can be found anymore. Indeed, by imposing explicit limits τ_{min} and τ_{max} on the minimum and maximum pheromone trails, one ensures that relative differences between pheromone trails cannot become too extreme. Therefore, the probability of choosing a vertex cannot become too small and stagnation situations are avoided. Furthermore, by initializing pheromone trails to τ_{max} at the beginning of the search, one ensures that during the first cycles the relative difference between pheromone trails is rather small (after i cycles, it is bounded by a ratio of ρ^i). Hence, exploration is emphasized at the beginning of the search.

The effectiveness of bounding pheromone trails, and initializing them to the upper bound, is demonstrated in [19] on the traveling salesman problem and the quadratic assignment problem. Also, it has been shown to be effective on constraint satisfaction problems [33].

To set appropriately τ_{min} and τ_{max}, we ran `Ant-Clique` on a "representative" subset of instances with different settings with $\tau_{max} \in [4, 10]$ and $\tau_{min} \in [0.001, 0.05]$. On average, the best results have been obtained with $\tau_{min} = 0.01$ and $\tau_{max} = 6$. Note that this setting is consistent with [19], in which it is shown that τ_{max} should be set to an estimate of the asymptotically maximum pheromone trail value, which is defined by $\delta_{avg}/(1-\rho)$ where δ_{avg} is the average quantity of pheromone that is laid on pheromonal components at each cycle. Indeed, in `Ant-Clique`, the quantity of pheromone laid at each cycle is equal to $1/(1+ \mid \mathcal{C}_{\text{best}} \mid - \mid \mathcal{C}_k \mid)$, where $\mathcal{C}_k$ is the largest clique built during the cycle, and $\mathcal{C}_{best}$ the largest clique built since the beginning of the run: this quantity is equal to 1 if $\mathcal{C}_k$ is the best clique found so far, and it is lower than 1 otherwise. This quantity varies from one cycle to another (depending on the size of the cliques

computed during the cycle), but also from an instance to another. On the set of DIMACS benchmark instances, the average gap of size beween $\mathcal{C}_k$ and $\mathcal{C}_{best}$ is close to 13, so that τ_{max} should be set to a value around $1/((1+13)\cdot(1-\rho))$ which is equal to 7.1 when $\rho = 0.99$.

3.2 Varying the number of ants

To emphasize diversification and avoid premature stagnation, one can also increase the number of ants so that more states are explored at each cycle. This parameter has been set experimentally, by running `Ant-Clique` with different values beween 10 and 50. On average, the best results have been obtained with $nbAnts = 30$. With lower values, solution quality is often decreased as the best clique constructed at each cycle usually is significantly smaller. With greater values, running time often increases while solution quality is not significantly improved as the best clique constructed at each cycle is not significantly better than with 30 ants. Experimental results and a deeper analysis about the influence of the number of ants on the solution process of `Edge-AC` can be found in [29].

3.3 Influence of α and ρ on the solution process

The two pheromonal parameters α and ρ have a great influence on the solution process. Indeed, diversification can be emphasized either by decreasing the value of the pheromone factor weight α —so that ants become less sensitive to pheromone trails— or by increasing the value of the pheromone persistence rate ρ —so that pheromone evaporates more slowly. When increasing the exploratory ability of ants in this way, one usually finds better solutions, but as a counterpart it takes longer time to find them. This is illustrated in Figure 2 on the `C500.9` instance of the DIMACS benchmark. Note that this instance, which has 500 vertices and 112332 edges, is a rather difficult one, for which `Ant-Clique` has difficulties in finding the maximum clique —which contains 57 vertices.

On this figure, one can first note that when $\alpha = 0$ and $\rho = 1$, the best constructed cliques are much smaller: in this case, pheromone is totally ignored and the resulting search process performs as a random one so that after 500 or so cycles, the size of the best clique nearly stops increasing, and hardly reaches 48 vertices. This shows that pheromone actually improves the solution process with respect to a pure random algorithm.

For both `Edge-AC` and `Vertex-AC`, we remark that α and ρ influence the solution process in a very similar way: when α increases or ρ decreases, ants are able to find better solutions quicker. However, after a thousand or so cycles, both versions find better solutions when α is set to 1 and ρ to 0.99 or 0.995 than when α is set to 2 or ρ to 0.98. Hence, the setting of α and ρ let us balance between two main tendencies. On the one hand, when limiting the influence of pheromone with a low pheromone factor weight and a high pheromone persistence rate, the quality of the final solution is better, but the time needed to converge on this value is also higher. On the other hand, when increasing the influence

Figure 2: Influence of pheromone on the solution process of `Edge-AC` (upper curves) and `Vertex-AC` (lower curves) for the `C500.9` graph: each curve plots the evolution of the size of the best constructed clique (average over 50 runs), when the number of cycles increases, for a given setting of α and ρ. The other parameters have been set to $nbAnts = 30$, $\tau_{min} = 0.01$, and $\tau_{max} = 6$.

11

of pheromone with a higher pheromone factor weight and a lower pheromone persistence rate, ants find better solutions during the first cycles, but after 500 or so cycles, they are no longer able to find better solutions.

When comparing `Edge-AC` (upper curves) with `Vertex-AC` (lower curves), one can note that after 2500 cycles, the best cliques found by `Edge-AC` are slightly larger than the best ones found by `Vertex-AC`. For example, when setting α to 1 and ρ to 0.99, the average size of the best constructed cliques is equal to 55.6 for `Edge-AC` and to 55.2 for `Vertex-AC`. Moreover, three runs of `Edge-AC` (over the fifty performed runs) have been able to find a clique of 57 vertices, whereas `Vertex-AC` only found cliques of 56 or less vertices.

However, if `Edge-AC` is able to find better cliques, it needs more cycles to converge towards them. For example, when setting α to 1 and ρ to 0.99, the average number of cycles needed to find the best clique is respectively equal to 923 for `Edge-AC` and 722 for `Vertex-AC`. Moreover, as discussed in Section 2.5, each cycle of `Edge-AC` takes a longer time to perform than `Vertex-AC`: on this instance, in one second of CPU time, `Edge-AC` and `Vertex-AC` respectively perform 104 and 188 cycles. Hence, to find their best cliques, `Edge-AC` and `Vertex-AC` respectively need 8.9 and 3.8 seconds on average.

As a conclusion, on instance `C500.9`, the two algorithms behave in a rather similar way at run time with respect to pheromone parameters, and a good compromise between solution quality and CPU time is reached when α is set to 1 and ρ to 0.99. With such a parameter setting, `Vertex-AC` is more than twice as fast as `Edge-AC` to find its best clique, but the best clique found by `Edge-AC` is slightly better, on average, than the one found by `Vertex-AC`, and `Edge-AC` has been able to find the best known solution for 6% of its runs, whereas `Vertex-AC` never reached it.

This study on instance `C500.9` allows us to determine a "good" parameter setting to solve this particular instance. However, one can hardly assert that the best parameters for `C500.9` are also the best ones for other instances. In particular, we have shown in [29] that the best parameters for another instance (`gen_400_P0.9_75`) were different: indeed this instance is "easier" than `C500.9` so that one has better choose parameters that favor a quick convergence such as $\alpha = 3$ and $\rho = 0.985$. In the rest of this paper, we have chosen to set parameters to the same values for all instances. Another possibility would have been to use an adaptive method, such as, e.g., the one described in [34]. This method aims at defining an automatic procedure for finding good parameters through statistically guided experimental evaluations. A racing method uses a pool of initial parameters and, as computation proceeds, gather evidences allowing to remove inferior candidates from the pool.

3.4 Measuring the diversity at run time

To provide an insight into `Ant-Clique`'s performance and to explicit the influence of pheromone on the capability of ants to explore the search space, we now propose to measure the diversity of the computed solutions at run time. Many diversity measures have been introduced for evolutionary approaches. Indeed,

Table 1: Evolution of the re-sampling ratio of `Edge-AC` and `Vertex-AC` for graph `C500.9`. Each row successively displays the setting of α and ρ, and the re-sampling ratio after 500, 1000, 1500, 2000, and 2500 cycles (average over 50 runs). The other parameters have been set to $nbAnts = 30$, $\tau_{min} = 0.01$, $\tau_{max} = 6$.

Re-sampling ratio for `Edge-AC`

Number of cycles:	500	1000	1500	2000	2500
$\alpha = 1, \rho = 0.995$	0.00	0.00	0.00	0.00	0.00
$\alpha = 1, \rho = 0.99$	0.00	0.00	0.00	0.00	0.00
$\alpha = 2, \rho = 0.99$	0.00	0.04	0.06	0.07	0.07
$\alpha = 2, \rho = 0.98$	0.06	0.10	0.12	0.13	0.13

Re-sampling ratio for `Vertex-AC`

Number of cycles:	500	1000	1500	2000	2500
$\alpha = 1, \rho = 0.995$	0.00	0.00	0.02	0.13	0.25
$\alpha = 1, \rho = 0.99$	0.00	0.07	0.26	0.39	0.48
$\alpha = 2, \rho = 0.99$	0.38	0.68	0.78	0.84	0.87
$\alpha = 2, \rho = 0.98$	0.58	0.78	0.85	0.88	0.91

maintaining population diversity is a key point to prevent from premature convergence and stagnation. Most commonly used diversity measures include the number of different fitness values, the number of different structural individuals, and distances between individuals [35].

To measure the diversification effort of `Ant-Clique` at run time, we propose in this paper to compute the re-sampling and the diversification ratio.

3.4.1 Re-sampling ratio

This measure is used, e.g., in [31, 36], in order to get insight into how "effective" algorithms are in sampling the search space: if we define $nbDiff$ as the number of unique candidate solutions generated by an algorithm over a whole run and $nbTot$ as the total number of generated candidate solutions, then the re-sampling ratio is defined as $(nbTot$ - $nbDiff)/nbTot$. Values close to 0 correspond to an "effective" search, i.e., not many duplicate candidate solutions are generated, whereas values close to 1 indicate a stagnation of the search process around a small set of solutions.

Table 1 provides an insight into `Ant-Clique`'s performance by means of this re-sampling ratio. This table shows that `Vertex-AC` is less "effective" than `Edge-AC` in sampling the search space as it often generates cliques that have already been previously generated. For example, when setting α to 1 and ρ to 0.99, 7% of the cliques computed by `Vertex-AC` during the first thousand of cycles had already been computed. This re-sampling ratio increases very quickly and reaches 48% at cycle 2500, i.e., nearly half of the cliques computed during

each run already had been computed before. As a comparison, with the same parameter setting, `Edge-AC` nearly never computes twice a same clique during a same run, so that it actually explores twice more states in the search space.

The re-sampling ratio allows one to quantify the size of the searched space, and shows that `Edge-AC` has a higher search capability than `Vertex-AC` (for the considered `C500.9` instance). However, the re-sampling ratio gives no information about the distribution of the computed cliques within the whole search space. Actually, a pure random search (nearly) never re-samples twice a same solution but, as it does not intensify the search around promising solutions, it usually is not able to find good solutions. Hence, to provide a complementary insight into `Ant-Clique`'s performance, we propose to compute a similarity ratio which indicates how much the computed cliques are similar, i.e., how much the search is intensified.

3.4.2 Similarity ratio

The similarity ratio corresponds to the pair-wise population diversity measure, introduced for genetic approaches, e.g., in [37, 38]. More precisely, we define the similarity ratio of a set of cliques S by the average number of vertices that are shared by any pair of cliques in S, divided by the average size of the cliques of S. Hence, this ratio is equal to one if all the cliques of S are identical, whereas it is equal to zero if the intersection of every pair of cliques of S is empty. Note that this ratio can be computed very quickly by maintaining an array `freq` such that, for every vertex $v_i \in V$, `freq[i]` is equal to the number of cliques of S which have selected vertex v_i. In this case, the similarity ratio of S is equal to

$$\frac{\sum_{v_i \in V}(\texttt{freq[i]} \cdot (\texttt{freq[i]} - 1))}{(\mid S \mid -1) \cdot \sum_{\mathcal{C}_k \in S} \mid \mathcal{C}_k \mid}$$

and it can be easily computed in an incremental way while constructing cliques (see [38] for more details).

Figure 3 plots the evolution of the similarity ratio of the cliques computed every 50 cycles, for graph `C500.9`, thus giving an information about the distribution of the set of cliques computed during these 50 cycles. For example, let us consider the curve plotting the evolution of the similarity ratio for `Edge-AC` when α is set to 1 and ρ to 0.99. The similarity increases from less than 10% at the beginning of the solution process to 45% after a thousand or so cycles. This shows that ants progressively focus on a sub-region of the search space, so that two cliques constructed after cycle 1000 share nearly half of their vertices on average. When considering this together with the fact that the re-sampling ratio is null, one can conclude that in this case `Edge-AC` reaches a good compromise between diversification —as it never re-computes twice a same solution— and intensification —as the similarity of the computed solutions is increased.

Figure 3 also shows that, when α increases or ρ decreases, the similarity ratio both increases sooner and rises more steeply. However, both for `Edge-AC` and `Vertex-AC`, the similarity ratio of all runs converges towards a same value,

Figure 3: Evolution of the similarity ratio for `Edge-AC` (upper curves) and `Vertex-AC` (lower curves) for graph `C500.9`: each curve plots the similarity ratio of the set of cliques constructed every 50 cycles (average over 50 runs), for a given setting of α and ρ. The other parameters have been set to $nbAnts = 30$, $\tau_{min} = 0.01$, $\tau_{max} = 6$.

whatever the setting of α and ρ: after two thousand or so cycles, the cliques computed by `Edge-AC` during every cycle share around 45% of their vertices whereas those computed by `Vertex-AC` share around 90% of their vertices.

The difference of diversification between `Edge-AC` and `Vertex-AC` may be explained by the choice made about their pheromonal components. Indeed, the considered `C500.9` instance has 500 vertices, so that in `Vertex-AC` ants may lay pheromone on 500 components, whereas in `Edge-AC` they may lay pheromone on $500 \cdot 499/2 = 124750$ components. Let us now consider the two cliques that are rewarded at the end of the first two cycles of a run. For both `Vertex-AC` and `Edge-AC`, these two cliques contain 44 vertices on average (see Fig. 2), and their similarity ratio is lower than 10% (see Fig. 3) so that they share 4 vertices on average. Under this hypothesis, after the first two cycles of `Vertex-AC`, 4 vertices —corresponding to 1% of the pheromonal components— have been rewarded twice, and 80 vertices —corresponding to 16% of the pheromonal components— have been rewarded once. As a comparison, under the same hypothesis, after the first two cycles of `Edge-AC`, 6 edges —corresponding to less than 0.005% of the pheromonal components— have been rewarded twice, and 940 edges —corresponding to less than 0.8% of the pheromonal components— have been rewarded once. This explains why the search of `Edge-AC` is much more diversified than the one of `Vertex-AC`, and therefore why `Edge-AC` needs more time to converge, but as a counterpart it often finds better solutions with respect to `Vertex-AC`.

4   Enhancing ACO with local search

Basically, local search searches for a locally optimal solution in the neighborhood of a given constructed solution. Local search may be combined with the ACO meta-heuristic in a very straightforward way: ants construct solutions exploiting pheromone trails, and local search improves their quality by iteratively performing local moves. Actually, the best-performing ACO algorithms for many combinatorial optimization problems are hybrid algorithms that combine probabilistic solution construction by a colony of ants with local search [22, 19, 28].

In this section, we study the benefit of integrating local search within `Ant-Clique`. The hybrid algorithm is derived from the algorithm of Figure 1 as follows: once each ant has constructed a clique, and before updating pheromone trails, we apply a local search procedure on the largest clique of the cycle until it becomes locally optimal[2]. Pheromone trails are then updated with respect to this locally optimal clique.

Various local search procedures may be used to improve cliques, e.g., [5, 6, 14]. However, as pointed out in [39], when choosing the local search procedure to use in a meta-heuristic such as evolutionary algorithms, iterated local search

---

[2]Local search could be applied to every computed clique (instead of applying it only to the best clique of the cycle). However, experiments showed us that this does not significantly improve solutions' quality, whereas it is much more time consuming.

or ACO, one has to find a trade-off between computation time and solution quality. In other words, one has to choose between a fast but not-so-good local search procedure or a slower but more drastic one.

For all experiments reported in this paper, we have considered the (2,1)-exchange procedure used in GRASP [5]. Given a clique $\mathcal{C}$, this local search procedure looks for three vertices v_i, v_j and v_k such that:
- v_i belongs to $\mathcal{C}$,
- v_j and v_k do not belong to $\mathcal{C}$,
- v_j and v_k are linked by an edge, and
- v_j and v_k are adjacent to every vertex of $\mathcal{C} - \{v_i\}$. Then, it replaces the vertex v_i by the two vertices v_j and v_k, thus increasing the clique size by one. This local search procedure is iterated until it reaches a locally optimal state that cannot be improved by such a (2,1)-exchange.

Figure 4 shows that this local search procedure actually enhances the performance of Ant-Clique when solving instance C500.9. In particular, when α is set to 0 and ρ to 1, i.e., when pheromone is not used, the local search procedure roughly increases the size of the constructed cliques by four, and this improvement in quality is constant during the whole run. When α is set to 1 and ρ to 0.99, so that pheromone actually influences the solution process, local search also improves the quality of the cliques constructed by Ant-Clique, but the improvement in quality is not constant during the run. At the beginning of the run, local search increases cliques by four vertices. However, after a thousand or so cycles, the improvement in quality is much smaller. Finally, at the end of the run the best clique found when combining Ant-Clique with local search is slightly larger than when local search is not used. Note that local search enhances the performances of the two variants of Ant-Clique in a very similar way: for Edge-AC (resp. Vertex-AC) the integration of local search increases the average quality of the best clique after 2500 cycles from 55.6 to 55.9 (resp. 55.2 to 55.4)

5 Experimental comparison of the four instantiations of Ant-Clique

In this section, we experimentally evaluate and compare the four proposed instantiations of Ant-Clique, i.e., Edge-AC, Vertex-AC, Edge-AC combined with local search (referred to as Edge-AC+LS) and Vertex-AC combined with local search (referred to as Vertex-AC+LS).

Experimental setup. All algorithms have been implemented in the C language and ran on a 1.9 GHz Pentium 4 processor, under Linux operating system.

In all experiments, we have set the number of ants to 30, τ_{min} to 0.01, τ_{max} to 6, α to 1, and ρ to 0.99, thus achieving a good compromise between solution quality and CPU time as discussed in Section 3.

The maximum number of cycles has been set to 5000. Indeed, for many

Figure 4: Evolution of the size of the best clique (average over 50 runs) for graph C500.9 with local search (*-AC + LS) or without local search (*-AC), and with pheromone (alpha=1, rho=0.99) or without pheromone (alpha=0, rho=1.00), for graph C500.9. The other parameters have been set to $nbAnts = 30$, $\tau_{min} = 0.01$, and $\tau_{max} = 6$. Upper curves display results for Edge-AC and lower curves for Vertex-AC.

benchmark instances, all algorithms have converged to the best solutions within the first three thousands of cycles. However, on larger instances —that have more than one thousand of vertices— `Edge-AC` may need more cycles to converge so that we have fixed the maximum number of cycles to 5000 for all instances.

Test suite. We consider 36 benchmark graphs provided by the DIMACS challenge on clique coloring and satisfiability[3].

- `C`n.p and `DSJC`n.p graphs are randomly generated graphs with n vertices and a density of 0.p, i.e., an edge is created between any pair of vertices with probability 0.p.

- `MANN_a27` and `MANN_a45` graphs are clique formulations of set covering formulations of Steiner Triple Problems and respectively have 378 and 1035 vertices, and a density of 0.99.

- `brock`n_m graphs have n vertices and a density between 0.496 and 0.74. These graphs contain large cliques hidden among a connected population of significantly smaller cliques.

- `gen`n_`p0.`p_m graphs have n vertices, a density of 0.p, and hidden cliques of known size m.

- `hamming`n_m graphs associate a vertex with each different n-bits word, so that `hamming8-4` and `hamming10-4` respectively have 256 and 1024 vertices. In these graphs, an edge exists between two vertices if and only if their words are at least distant of a Hamming distance m, so that the density respectively is 0.639 and 0.829.

- `Keller` graphs are based on a representation of Keller's conjecture on the tiling of space using hypercubes. `keller` 4, 5 and 6 graphs respectively have 171, 776 and 3361 vertices, and a density of 0.649, 0.751, and 0.818.

- `p_hat`n_m are random graphs having n vertices, and a density between 0.244 and 0.506. These graphs are tuned to have a wider vertex degree spectrum.

We do not report results for the `MANN_a81` instance, which has 3321 vertices and a maximum clique of 1098 vertices: for this instance, `Vertex-AC` and `Edge-AC` respectively spend 8800 and 11600 seconds of CPU time to perform 5000 cycles. Note that this prohibitive time is mainly due to the size of the computed cliques. Indeed, the `keller6` graph, that has 3361 vertices but a maximum clique of 59 vertices, is processed in a few minutes by both `Vertex-AC` and `Edge-AC`.

---

[3]This benchmark is available at `http://dimacs.rutgers.edu/`.

Table 2: Comparison of ACO algorithms by means of solutions' quality. Each row successively displays the instance name, the size of its maximum clique, and the results obtained by `Vertex-AC`, `Edge-AC`, `Vertex-AC+LS`, and `Edge-AC+LS` (best and average solution found over 50 runs; standard deviation in brackets).

		Size of the best found clique							
		Vertex-AC		Edge-AC		Vertex-AC+LS		Edge-AC+LS	
Graph	$\omega(G)$	Max	Avg(Stdv)	Max	Avg(Stdv)	Max	Avg(Stdv)	Max	Avg(Stdv)
C125.9	34	34	34.0 (0.0)	34	34.0 (0.0)	34	34.0 (0.0)	34	34.0 (0.0)
C250.9	44	44	43.9 (0.3)	44	44.0 (0.0)	44	44.0 (0.1)	44	44.0 (0.0)
C500.9	≥57	56	55.2 (0.8)	57	55.6 (0.8)	56	55.4 (0.8)	57	55.9 (0.6)
C1000.9	≥68	67	65.3 (1.0)	67	66.0 (0.8)	67	65.7 (0.6)	68	66.2 (0.7)
C2000.9	≥78	76	73.4 (1.1)	76	74.1 (1.3)	77	74.5 (1.1)	78	74.3 (1.4)
DSJC500.5	14	13	13.0 (0.0)	13	13.0 (0.0)	13	13.0 (0.0)	13	13.0 (0.0)
DSJC1000.5	15	15	14.1 (0.3)	15	14.1 (0.4)	15	14.2 (0.4)	15	14.3 (0.4)
C2000.5	≥16	16	14.9 (0.4)	16	15.1 (0.3)	16	15.1 (0.3)	16	15.3 (0.5)
C4000.5	≥18	17	15.9 (0.4)	16	15.8 (0.4)	17	16.2 (0.4)	18	16.8 (0.6)
MANN_a27	126	126	125.5 (0.5)	126	126.0 (0.2)	126	125.8 (0.4)	126	126.0 (0.0)
MANN_a45	345	344	342.8 (0.8)	344	343.3 (0.6)	344	342.6 (0.7)	344	342.9 (0.6)
brock200_2	12	12	11.9 (0.2)	12	12.0 (0.0)	12	11.9 (0.4)	12	12.0 (0.0)
brock200_4	17	17	16.1 (0.3)	17	16.8 (0.4)	17	16.1 (0.3)	17	16.8 (0.4)
brock400_2	29	25	24.5 (0.5)	29	25.0 (0.7)	25	24.7 (0.5)	25	24.8 (0.4)
brock400_4	33	25	24.0 (0.1)	33	25.1 (2.7)	25	24.2 (0.4)	33	27.1 (4.0)
brock800_2	24	21	20.0 (0.5)	21	19.8 (0.5)	21	20.4 (0.5)	24	20.1 (0.6)
brock800_4	26	21	19.8 (0.6)	26	19.9 (1.0)	21	20.2 (0.4)	26	20.0 (0.8)
gen200_p0.9_44	44	44	41.4 (1.9)	44	43.7 (1.1)	44	43.3 (1.5)	44	44.0 (0.0)
gen200_p0.9_55	55	55	55.0 (0.0)	55	55.0 (0.0)	55	55.0 (0.0)	55	55.0 (0.0)
gen400_p0.9_55	55	52	51.2 (0.5)	53	51.9 (0.5)	52	51.3 (0.5)	53	52.2 (0.4)
gen400_p0.9_65	65	65	65.0 (0.0)	65	65.0 (0.0)	65	65.0 (0.0)	65	65.0 (0.0)
gen400_p0.9_75	75	75	75.0 (0.0)	75	75.0 (0.0)	75	75.0 (0.0)	75	75.0 (0.0)
hamming8_4	16	16	16.0 (0.0)	16	16.0 (0.0)	16	16.0 (0.0)	16	16.0 (0.0)
hamming10_4	40	40	38.0 (1.5)	40	38.6 (1.2)	39	38.7 (0.6)	40	39.3 (0.9)
keller4	11	11	11.0 (0.0)	11	11.0 (0.0)	11	11.0 (0.0)	11	11.0 (0.0)
keller5	27	27	26.7 (0.5)	27	26.9 (0.2)	27	26.9 (0.3)	27	27.0 (0.0)
keller6	≥59	55	50.8 (1.9)	59	53.1 (1.7)	55	51.5 (1.5)	57	55.1 (1.3)
p_hat300_1	8	8	8.0 (0.0)	8	8.0 (0.0)	8	8.0 (0.0)	8	8.0 (0.0)
p_hat300_2	25	25	25.0 (0.0)	25	25.0 (0.0)	25	25.0 (0.0)	25	25.0 (0.0)
p_hat300_3	36	36	35.9 (0.5)	36	36.0 (0.1)	36	36.0 (0.3)	36	36.0 (0.0)
p_hat700_1	11	11	10.8 (0.4)	11	11.0 (0.1)	11	10.9 (0.3)	11	11.0 (0.1)
p_hat700_2	44	44	44.0 (0.0)	44	44.0 (0.0)	44	44.0 (0.0)	44	44.0 (0.0)
p_hat700_3	≥62	62	62.0 (0.0)	62	62.0 (0.0)	62	62.0 (0.0)	62	62.0 (0.0)
p_hat1500_1	12	12	11.0 (0.2)	12	11.1 (0.3)	12	11.2 (0.4)	12	11.1 (0.2)
p_hat1500_2	≥65	65	64.9 (0.2)	65	64.9 (0.2)	65	65.0 (0.0)	65	65.0 (0.0)
p_hat1500_3	≥94	94	93.1 (0.2)	94	93.9 (0.4)	94	93.6 (0.5)	94	94.0 (0.0)

Table 3: Comparison of ACO algorithms by means of CPU time. Each row successively displays the number of cycles (average over 50 runs) and the CPU time (average and standard deviation over 50 runs) for `Vertex-AC`, `Edge-AC`, `Vertex-AC+LS`, and `Edge-AC+LS`.

	Number of cycles and time to find the best clique							
	Vertex-AC		Edge-AC		Vertex-AC+LS		Edge-AC+LS	
Graph	Cycles	Time(Stdv)	Cycles	Time(Stdv)	Cycles	Time(Stdv)	Cycles	Time(Stdv)
C125.9	60	0.1 (0.0)	126	0.2 (0.1)	14	0.0 (0.0)	23	0.0 (0.0)
C250.9	359	0.8 (0.7)	473	1.7 (0.3)	172	0.5 (0.1)	239	1.0 (0.3)
C500.9	722	3.8 (1.9)	923	8.9 (4.0)	477	4.6 (2.7)	671	8.6 (4.7)
C1000.9	1219	13.2 (5.8)	2359	55.0 (21.9)	832	23.4 (8.5)	1242	49.8 (26.6)
C2000.9	1770	41.3 (11.0)	3278	214.4 (49.8)	1427	112.4 (16.8)	2067	238.7 (98.3)
DSJC500.5	588	0.5 (0.2)	832	2.6 (1.8)	249	0.7 (0.5)	285	1.4 (1.3)
DSJC1000.5	820	1.4 (1.0)	1017	9.6 (6.4)	561	3.8 (4.0)	567	7.8 (8.8)
C2000.5	957	3.3 (3.0)	1062	30.4 (16.8)	312	5.9 (5.7)	957	40.6 (56.0)
C4000.5	927	6.0 (5.1)	1116	108.3 (73.8)	445	22.4 (33.1)	1915	257.6 (190.2)
MANN_a27	616	11.5 (3.5)	2274	61.7 (25.3)	574	11.1 (4.9)	1824	44.8 (19.3)
MANN_a45	1771	271.3 (77.0)	4360	877.4 (69.1)	1498	239.2 (56.9)	3639	749.4 (132.5)
brock200_2	115	0.0 (0.0)	127	0.1 (0.1)	104	0.1 (0.1)	115	0.1 (0.1)
brock200_4	156	0.1 (0.0)	1627	1.3 (1.2)	69	0.1 (0.0)	1356	1.7 (1.9)
brock400_2	650	1.0 (0.6)	1004	3.5 (1.9)	424	1.4 (0.6)	720	3.8 (3.5)
brock400_4	498	0.8 (0.2)	977	3.4 (3.3)	254	0.8 (0.4)	1141	5.7 (6.3)
brock800_2	1145	2.6 (1.1)	1238	9.1 (6.1)	881	6.3 (3.6)	959	11.6 (10.5)
brock800_4	1173	2.7 (1.0)	1288	9.8 (7.1)	858	6.1 (2.5)	906	11.1 (10.7)
gen200_p0.9_44	297	0.6 (0.2)	338	0.9 (0.2)	136	0.3 (0.1)	165	0.5 (0.1)
gen200_p0.9_55	102	0.2 (0.0)	130	0.3 (0.1)	57	0.2 (0.0)	100	0.3 (0.1)
gen400_p0.9_55	475	2.0 (1.2)	1634	11.5 (6.4)	270	1.8 (0.9)	760	6.7 (3.7)
gen400_p0.9_65	337	1.4 (0.2)	391	2.7 (0.3)	206	1.5 (0.1)	255	2.3 (0.3)
gen400_p0.9_75	251	1.0 (0.1)	293	2.1 (0.2)	153	1.2 (0.1)	198	2.1 (0.2)
hamming8_4	34	0.0 (0.0)	42	0.1 (0.0)	49	0.1 (0.1)	42	0.1 (0.1)
hamming10_4	2308	13.7 (1.9)	1474	28.0 (10.5)	1204	26.2 (17.3)	865	29.3 (16.3)
keller4	2	0.0 (0.0)	2	0.0 (0.0)	0	0.0 (0.0)	0	0.0 (0.0)
keller5	1652	4.9 (1.2)	1136	10.8 (7.9)	979	9.4 (6.0)	830	12.3 (9.7)
keller6	2514	55.3 (15.2)	3034	308.8 (111.8)	1657	206.8 (145.3)	2617	549.2 (250.6)
p_hat300_1	40	0.0 (0.0)	46	0.0 (0.0)	18	0.0 (0.0)	20	0.0 (0.0)
p_hat300_2	113	0.1 (0.0)	206	0.3 (0.1)	26	0.1 (0.0)	54	0.2 (0.1)
p_hat300_3	281	0.5 (0.5)	457	1.3 (0.3)	110	0.3 (0.1)	176	0.7 (0.2)
p_hat700_1	379	0.2 (0.1)	624	2.6 (1.6)	253	0.7 (0.3)	391	2.4 (2.0)
p_hat700_2	297	0.7 (0.1)	445	3.2 (0.6)	128	2.0 (0.6)	227	4.5 (1.0)
p_hat700_3	575	3.0 (2.2)	878	9.8 (3.2)	220	3.8 (1.0)	333	7.8 (1.4)
p_hat1500_1	391	0.4 (0.3)	662	10.5 (8.5)	318	3.0 (4.4)	438	9.7 (17.0)
p_hat1500_2	499	3.3 (1.0)	801	20.1 (4.9)	238	17.3 (2.7)	379	35.2 (3.9)
p_hat1500_3	581	8.6 (6.5)	2199	74.3 (41.6)	551	37.0 (34.4)	528	54.9 (8.8)

Comparison of solutions' quality. Table 2 compares the four proposed ACO instantiations by means of solutions' quality. On this table, one can first note that `Edge-AC` outperforms `Vertex-AC`: when considering average results over the 50 runs, `Edge-AC` has found larger cliques than `Vertex-AC` for 21 instances, whereas it has found smaller cliques for 2 instances only. Moreover, `Edge-AC` has been able to find the best known solution for 29 instances, whereas `Vertex-AC` has found it for 24 instances.

This table also shows that local search improves the solution process of both algorithms for many instances. For `Vertex-AC`, the integration of local search improves average results for 22 instances, whereas it deteriorates them for 1 instance, and for `Edge-AC`, the integration of local search improves average results for 16 instances, whereas it deteriorates them for 2 instances. Note that the benefit of integrating local search within `Ant-Clique` varies with respect to the considered classes of instances. In particular, local search does not improve solutions' quality for `MANN` instances, whereas it significantly improves it on hard random `C*.*` instances, so that `Edge-AC+LS` has been able to find the best known solutions for all `C*.*` instances.

Hence, `Edge-AC+LS` is the best performing of the four proposed ACO instantiations, for a majority of instances, and it has been able to find the best known solution for 31 instances, over the 36 considered instances. However, for 2 instances (`brock400_2` and `keller6`), `Edge-AC` has found the best known solution whereas `Edge-AC+LS` did not.

Finally, this comparison on a large set of instances has shown us that solutions' quality depends on the considered classes of instances. In particular, the four considered ACO instantiations obtained very good average results on `gen*_p*_*` and `p_hat*_*` instances, being able to find the best known solution for nearly all the runs on most of these instances. On the other side, average results are rather far from optimality on `brock*_*` instances. Actually, as pointed out in [40], `brock*_*` graphs were created using trap mechanisms destined to fool greedy algorithms and guide them towards deceptively easy solutions while bigger cliques remain hidden. As a consequence, if some runs have been able to find these hidden cliques, most runs have been trapped in smaller cliques.

Comparison of CPU time. Table 3 compares the four proposed ACO instantiations by means of CPU time. This table first shows that, for each considered ACO instantiation, the number of cycles —and therefore the CPU time— needed to find the best solution mainly depends on the number of vertices and the connectivity of the graph. For example, `Edge-AC` respectively performs, on average, 473, 923, 2359, and 3278 cycles to solve instances `C250.9`, `C500.9`, `C1000.9`, and `C2000.9` that respectively have 250, 500, 1000 and 2000 vertices and a connectivity close to 0.9, whereas it performs 1062 cycles to solve instance `C2000.5` that has 2000 vertices and a connectivity of 0.5.

Table 3 also shows that `Vertex-AC` nearly always performs less cycles than `Edge-AC` and, on average for all instances, it needs 1.8 times less cycles to converge towards its best solution. When considering CPU times, the difference
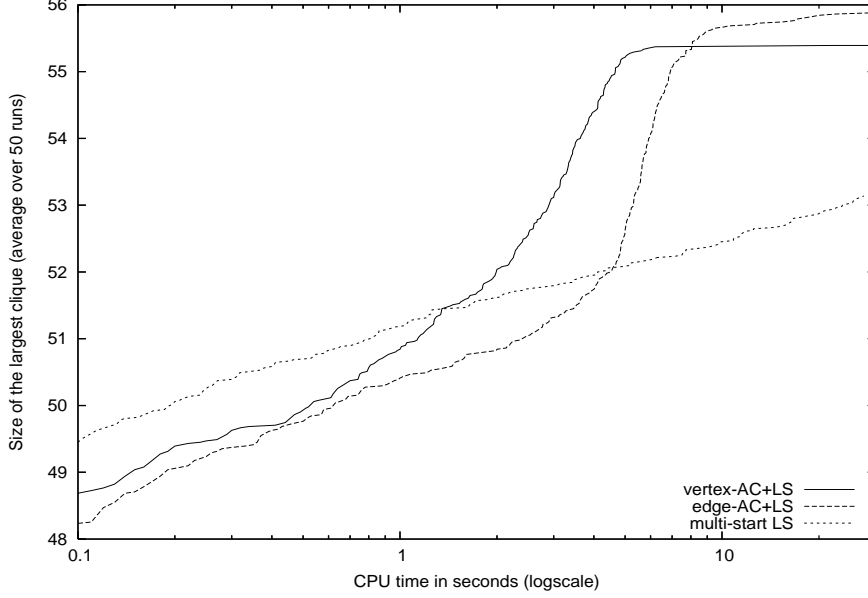
Figure 5: Comparison of `Edge-AC+LS`, `Vertex-AC+LS`, and multi-start local search on graph `C500.9`. Parameters have been set to $\alpha = 1$, $\rho = 0.99$, $nbAnts = 30$, $\tau_{min} = 0.01$, and $\tau_{max} = 6$ for `Edge-AC+LS` and `Vertex-AC+LS`.

becomes more important as `Edge-AC` needs more time to perform one cycle than `Vertex-AC`. Hence, `Vertex-AC` is from 1.5 to 26 times as fast as `Edge-AC`, and it is 5.7 times as fast on average for all instances.

Finally, one can remark that the integration of local search always decreases the number of cycles: on average, both `Vertex-AC+LS` and `Edge-AC+LS` perform 1.8 times less cycles than `Vertex-AC` and `Edge-AC` respectively. However, as local search is time consuming, CPU times are rather comparable.

Discussion. This comparison on 36 benchmark graphs showed us that integrating local search within `Ant-Clique` nearly always improves solutions' quality without increasing CPU times, so that one clearly has better use `Vertex-AC+LS` or `Edge-AC+LS` instead of `Vertex-AC` or `Edge-AC`.

However, to choose between `Vertex-AC+LS` and `Edge-AC+LS`, one has to consider the CPU time available for the solution process. Indeed, `Edge-AC+LS` usually finds better solutions than `Vertex-AC+LS` at the end of its solution process, but it is also more time consuming. As a consequence, if one has to find a solution within a short time limit, one has better use `Vertex-AC+LS`, whereas for larger time limits, or when there is no time limit, one has better use `Edge-AC+LS`. This is illustrated in Figure 5 on graph `C500.9`: for time limits smaller than 8 seconds, `Vertex-AC+LS` finds larger cliques than `Edge-AC+LS`, whereas for larger time limits `Edge-AC+LS` finds larger cliques than `Vertex-AC+LS`.

Figure 5 also compares `Vertex-AC+LS` and `Edge-AC+LS` with a multi-start

23

local search procedure called `multi-start LS`. This multi-start local search procedure iterates on the two following steps:

1. randomly build a maximal clique,

2. apply the local search procedure described in Section 4 on this maximal clique

Hence, `multi-start LS` corresponds to `Ant-Clique+LS` when pheromone is not used and the number of ants is set to 1 (so that local search is applied to every constructed clique). Figure 5 shows us that during the first second of CPU time `multi-start LS` finds better solutions than `Ant-Clique+LS`. Indeed, `Ant-Clique+LS` spends time to manage pheromone whereas this pheromone starts influencing ants only after a few hundreds or so cycles. Hence, `Vertex-AC+LS` (resp.`Edge-AC+LS`) finds better solutions than `multi-start LS` only after one second (resp. five seconds) of CPU time.

6 Experimental comparison with other heuristic approaches

We now compare `Edge-AC+LS`, which is our best performing ACO algorithm by means of solution quality on a majority of instances, with three recent and representative algorithms, i.e., `RLS`, `DAGS`, and `GLS`.

RLS (Reactive Local Search) [6] is based on a tabu local search heuristic. Starting from an empty clique, `RLS` iteratively moves in the search space composed of all cliques by adding/removing one vertex to/from the current clique. A tabu list is used to memorize the T last moves and, at each step, `RLS` greedily selects a move that is not prohibited by the tabu list. The key point is that the length T of the tabu list is dynamically updated with respect to the need for diversification. `RLS` appears to be the best heuristic algorithm for the maximum clique problem we are aware of, for a majority of DIMACS benchmark instances.

Table 4 displays results reported in [6] for the "basic" version of RLS which obtained the best results. CPU times have been multiplied by 0.027, corresponding to the ratio of computers' speed with respect to the Spec benchmark on floating point units.

DAGS [8] is a two-phase procedure: in a first phase, a greedy procedure combined with "swap" local moves is applied, starting from each node of the graph; in a second phase, nodes are scored with respect to the number of times they have been selected during the first phase, and an adaptive greedy algorithm is repeatedly started to build cliques around the nodes with the least scores, in order to diversify the search towards less explored areas.

Experiments reported in [8] show that the first phase is able to find best known solutions for many instances. For harder instances, the second phase

24

improves solutions' quality in many cases. This improvement in quality is dramatic on the set of `brock` instances, that are known to be very difficult for greedy approaches. For these instances, `DAGS` outperforms `RLS` performance.

Table 4 displays results reported in [8]. Note that the second phase of `DAGS` has been performed only for the harder instances, that have not been solved during the first phase. Hence, we specify in Table 4 if these results have been obtained after the first or the second phase. We multiplied CPU times by 0.84, corresponding to the ratio of computers' speed with respect to the Spec benchmark on floating point units. Note also that CPU times reported for DAGS correspond to the total time spent for the whole solving process, and not the time spent to find the best solution like for the three other considered algorithms.

GLS [14] combines a genetic algorithm with local search. We consider it for comparison, though it does not outperform `RLS` nor `DAGS`, because it presents some similarities with ACO: both approaches use a bio-inspired metaphor to intensify the search towards the most "promising" areas with respect to previously computed solutions. `GLS` generates successive populations of maximal cliques from an initial one by repeatedly selecting two parent cliques from the current population, recombining them to generate two children cliques, applying local search on children to obtain maximal cliques, and adding to the new population the best two cliques of parents and children. `GLS` can be instantiated to different algorithms by modifying its parameters. In particular, [14] compares results obtained by the three following instantiations of `GLS`: `GENE` performs genetic local search; `ITER` performs iterated local search, starting from one random point; and `MULT` performs multi-start local search, starting from a new random point at each time.

Table 4 displays, for each considered instance, the results obtained by the `GLS` algorithm (over `GENE`, `ITER`, and `MULT`) that obtained the best average results. We multiplied CPU times by 0.037, corresponding to the ratio of computers' speed with respect to the Spec benchmark on floating point units.

By means of solutions' quality, the results of table 4 are summarized in table 5. When considering the best found results, `Edge-AC+LS` is competitive with both `RLS` and `DAGS`, being able to find better solutions than `RLS` on two `brock` instances, and better solutions than `DAGS` on two `C*.9` instances. However, when considering average results `RLS` outperforms `Edge-AC+LS` on 16 instances. The comparison with `DAGS` on average results depends on the considered instances. In particular, `DAGS` outperforms `Edge-AC+LS` on `brock` instances, whereas `Edge-AC+LS` outperforms `DAGS` on `gen` instances.

This table also shows that `Edge-AC+LS` outperforms `GLS` both with respect to best and average results. Hence, for this problem ACO is better suited than evolutionary computation for guiding the search towards promising areas.

When considering CPU times reported in Table 4, one can note that `Edge-`

Table 4: Comparison of `Edge-AC+LS` with other heuristic approaches. Each row successively displays the results obtained by `Edge-AC+LS`, `RLS`, `DAGS` (followed by (1) or (2) if they have been obtained after the first or the second phase), and `GLS` (followed by (G), (I), or (M) if they have been obtained by `GENE`, `ITER`, or `MULT`). For each approach, the table reports the best and average solution found (over 50 runs for `Edge-AC+LS`, 100 runs for `RLS`, 20 runs `DAGS`, and 10 runs for `GLS`), and the estimated CPU time w.r.t. the Spec floating points benchmark.

	Edge-AC+LS			RLS			DAGS				GLS			
Graph	Max	Avg	Time	Max	Avg	Time	Max	Avg	Time		Max	Avg	Time	
C125.9	34	34.0	0.0	34	34.0	0.0	34	34.0	0.2	(1)	34	34.0	0.0	(I)
C250.9	44	44.0	1.0	44	44.0	0.0	44	44.0	0.7	(1)	44	43.0	0.1	(I)
C500.9	57	55.9	8.6	57	57.0	0.0	56	55.8	13.5	(2)	55	52.7	0.1	(I)
C1000.9	68	66.2	49.8	68	68.0	1.1	68	65.9	148.2	(2)	66	61.6	0.5	(G)
C2000.9	78	74.3	238.7	78	77.6	22.1	76	75.4	1824.0	(2)	70	68.7	0.9	(I)
DSJC500.5	13	13.0	1.4	13	13.0	0.0	-	-	-	-	13	12.2	0.1	(G)
DSJC1000.5	15	14.3	7.8	15	15.0	0.2	-	-	-	-	14	13.5	0.1	(I)
C2000.5	16	15.3	40.6	16	16.0	0.3	16	15.9	24.8	(1)	15	14.2	0.1	(I)
C4000.5	18	16.8	257.6	18	18.0	58.7	18	17.5	3229.0	(2)	16	15.6	0.6	(I)
MANN_a27	126	126.0	44.8	126	126.0	0.1	126	126.0	6.1	(1)	126	126.0	0.6	(I)
MANN_a45	344	342.9	749.4	345	343.6	10.7	344	343.9	1921.0	(2)	345	343.1	2.0	(I)
brock200_2	12	12.0	0.1	12	12.0	0.3	12	12.0	0.1	(2)	12	12.0	0.1	(M)
brock200_4	17	16.8	1.7	17	17.0	0.5	17	16.8	0.3	(2)	17	15.7	0.1	(M)
brock400_2	25	24.8	3.8	29	26.0	1.1	29	28.1	2.8	(2)	25	23.2	0.1	(I)
brock400_4	33	27.1	5.7	33	32.4	2.9	33	33.0	2.8	(2)	25	23.6	0.0	(G)
brock800_2	24	20.1	11.6	21	21.0	0.1	24	20.8	16.8	(2)	20	19.3	0.2	(G)
brock800_4	26	20.0	11.1	21	21.0	0.2	26	22.6	16.9	(2)	20	19.0	0.1	(I)
gen200_p0.9_44	44	44.0	0.5	44	44.0	0.0	44	41.1	0.9	(2)	44	39.7	0.1	(G)
gen200_p0.9_55	55	55.0	0.3	55	55.0	0.0	55	55.0	0.4	(1)	55	50.8	0.1	(G)
gen400_p0.9_55	53	52.2	6.7	55	55.0	0.0	53	51.8	7.2	(2)	55	49.7	0.1	(G)
gen400_p0.9_65	65	65.0	2.3	65	65.0	0.0	65	55.4	7.3	(2)	65	53.7	0.2	(G)
gen400_p0.9_75	75	75.0	2.1	75	75.0	0.0	75	55.2	7.8	(2)	75	62.7	0.2	(I)
hamming8_4	16	16.0	0.1	16	16.0	0.0	-	-	-	-	16	16.0	0.0	(G)
hamming10_4	40	39.3	29.3	40	40.0	0.0	40	40.0	12.8	(1)	40	38.2	0.2	(I)
keller4	11	11.0	0.0	11	11.0	0.0	-	-	-	-	11	11.0	0.0	(G)
keller5	27	27.0	12.3	27	27.0	0.0	27	27.0	4.3	(1)	27	26.3	0.2	(I)
keller6	57	55.1	549.2	59	59.0	5.1	57	56.4	12326.0	(2)	56	52.7	1.3	(I)
p_hat300_1	8	8.0	0.0	8	8.0	0.0	8	8.0	0.1	(1)	8	8	0.0	(G)
p_hat300_2	25	25.0	0.2	25	25.0	0.0	25	25.0	0.5	(1)	25	25.0	0.0	(I)
p_hat300_3	36	36.0	0.7	36	36.0	0.0	36	36.0	0.8	(1)	36	35.1	0.1	(I)
p_hat700_1	11	11.0	2.4	11	11.0	0.0	11	11.0	0.7	(1)	11	9.9	0.1	(I)
p_hat700_2	44	44.0	4.5	44	44.0	0.0	44	44.0	5.5	(1)	44	43.6	0.0	(I)
p_hat700_3	62	62.0	7.8	62	62.0	0.0	62	62.0	8.5	(1)	62	61.8	0.2	(I)
p_hat1500_1	12	11.1	9.7	12	12.0	0.8	12	11.7	31.1	(2)	11	10.8	0.5	(G)
p_hat1500_2	65	65.0	35.2	65	65.0	0.0	65	65.0	47.7	(1)	65	63.9	0.5	(I)
p_hat1500_3	94	94.0	54.9	94	94.0	0.0	94	94.0	82.2	(1)	94	93.0	0.3	(I)

Table 5: Number of times `Edge-AC+LS` has found larger (+), equal (=), and smaller (−) cliques than `RLS`, `DAGS` and `GLS`.

	Comparison of `Edge-AC+LS` with								
	RLS			DAGS			GLS		
	+	=	−	+	=	−	+	=	−
Best clique	2	30	4	2	29	1	11	23	2
Average size	0	20	16	6	15	11	28	7	1

`AC+LS` is an order slower than `RLS`, which itself is an order slower than `GLS`. However, `Edge-AC+LS` is rather comparable with `DAGS`: it is quicker than `DAGS` for 24 instances, and it is slower for 8 instances.

7   Conclusion

We have described a generic ACO algorithm for searching for maximum cliques, and two different instantiations for it. These two instantiations differ in the choice of their pheromonal components. A main motivation was to answer the following question: should we lay pheromone on the vertices or on the edges of the graph?

Experiments have shown us that both algorithms are able to find optimal solutions on many benchmark instances, showing that ACO is actually able to guide the search towards promising areas. However, when comparing the diversification capability of the two algorithms by means of re-sampling and similarity ratio, we have noticed that the search is more diversified when pheromone is laid on edges so that better solutions are found on a wide majority of benchmark instances, but as a counterpart, more time is needed to converge. Also, experiments have shown us that the integration of local search techniques improves the solution process and makes ACO competitive with state-of-the-art heuristic approaches, though it is more time consuming.

We believe that this comparison of two ACO models for the maximum clique problem could be useful to solve other similar problems. Indeed, for many combinatorial optimization problems such as multi-dimensional knapsack problems or generalized assignment problems, the goal is to find, given an initial set of objects, the best subset with respect to some objective function. To solve this kind of problems with ACO, one has to choose between laying pheromone on the objects to choose, or on edges linking the objects to choose. In the first case, one will increase the desirability of choosing each rewarded object independently from the others, whereas in the second case, one will increase the desirability of choosing together two objects. Hence, further work will concern a generalization of this comparative study.

References

[1] R.M. Karp. Reducibility among combinatorial problems. *Complexity of Computer Computations*, pages 85–104, 1972.

[2] I. Bomze, M. Budinich, P.M. Pardalos, and M. Pelillo. The maximum clique problem. In D.-Z. Du and P. M. Pardalos, editors, *Handbook of Combinatorial Optimization*, volume 4, pages 1–74. Springer, 1999.

[3] D.S. Johnson. Approximation algorithms for combinatorial problems. *Journal of Computer Science*, 9:256–278, 1974.

[4] T.A. Feo, M.G.C. Resende, and S.H. Smith. A greedy randomized adaptive search procedure for maximum independent set. *Operations Research*, 42:860–878, 1994.

[5] J. Abello, P.M. Pardalos, and M.G.C. Resende. On maximum clique problems in very large graphs. In J.M. Abello, editor, *External Memory Algorithms*, volume 50 of *DIMACS Series in Discrete Mathematics and Theoretical Computer Science*, pages 119–130. American Mathematical Society, Boston, MA, USA, 1999.

[6] R. Battiti and M. Protasi. Reactive local search for the maximum clique problem. *Algorithmica*, 29(4):610–637, 2001.

[7] A. Jagota and L.A. Sanchis. Adaptive, restart, randomized greedy heuristics for maximum clique. *Journal of Heuristics*, 7(6):565–585, 2001.

[8] A. Grosso, M. Locatelli, and F. Della Croce. Combining swaps and node weights in an adaptive greedy approach for the maximum clique problem. *Journal of Heuristics*, 10(2):135–152, 2004.

[9] E.H.L. Aarts and J.H.M. Korst. *Simulated annealing and Boltzmann machines: a stochastic approach to combinatorial optimization and neural computing*. John Wiley & Sons, Chichester, U.K., 1989.

[10] S. Homer and M. Peinado. Experiments with polynomial-time clique approximation algorithms on very large graphs. In D.S. Johnson and M.A. Trick, editors, *Cliques, Coloring, and Satisfiability*, volume 26 of *DIMACS Series in Discrete Mathematics and Theoretical Computer Science*, pages 147–168. American Mathematical Society, Boston, MA, USA, 1996.

[11] C. Friden, A. Hertz, and D. de Werra. STABULUS: A technique for finding stable sets in large graphs with tabu search. *Computing*, 42(1):35–44, 1989.

[12] F. Glover and M. Laguna. Tabu search. In C.R. Reeves, editor, *Modern Heuristics Techniques for Combinatorial Problems*, pages 70–141. Blackwell Scientific Publishing, Oxford, UK, 1993.

[13] M. Gendreau, P. Soriano, and L. Salvail. Solving the maximum clique problem using a tabu search approach. *Annals of Operations Research*, 41(4):385–403, 1993.

[14] E. Marchiori. Genetic, iterated and multistart local search for the maximum clique problem. In S. Cagnoni, J. Gottlieb, E. Hart, M. Middendorf, and G.R. Raidl, editors, *Applications of Evolutionary Computing, Proceedings of EvoWorkshops 2002: EvoCOP, EvoIASP, EvoSTim*, volume 2279 of LNCS, pages 112–121. Springer-Verlag, 2002.

[15] M. Dorigo and G. Di Caro. The Ant Colony Optimization meta-heuristic. In D. Corne, M. Dorigo, and F. Glover, editors, *New Ideas in Optimization*, pages 11–32. McGraw Hill, UK, 1999.

[16] M. Dorigo, G. Di Caro, and L.M. Gambardella. Ant algorithms for discrete optimization. *Artificial Life*, 5(2):137–172, 1999.

[17] T. Jones and S. Forrest. Fitness distance correlation as a measure of problem difficulty for genetic algorithms. In L. Eshelman, editor, *Proceedings of the Sixth International Conference on Genetic Algorithms*, pages 184–192, San Francisco, CA, 1995. Morgan Kaufmann.

[18] P. Merz and B. Freisleben. Fitness landscapes and memetic algorithm design. In D. Corne, M. Dorigo, and F. Glover, editors, *New Ideas in Optimization*, pages 245–260. McGraw Hill, UK, 1999.

[19] T. Stützle and H.H. Hoos. $\mathcal{MAX} - \mathcal{MIN}$ Ant System. *Journal of Future Generation Computer Systems, special isuue on Ant Algorithms*, 16:889–914, 2000.

[20] M. Dorigo. *Optimization, Learning and Natural Algorithms* (in Italian). PhD thesis, Dipartimento di Elettronica, Politecnico di Milano, Italy, 1992.

[21] M. Dorigo, V. Maniezzo, and A. Colorni. Ant System: Optimization by a colony of cooperating agents. *IEEE Transactions on Systems, Man, and Cybernetics – Part B: Cybernetics*, 26(1):29–41, 1996.

[22] M. Dorigo and L.M. Gambardella. Ant colony system: A cooperative learning approach to the traveling salesman problem. *IEEE Transactions on Evolutionary Computation*, 1(1):53–66, 1997.

[23] L.M. Gambardella, E. Taillard, and M. Dorigo. Ant colonies for the quadratic assignment problem. *Journal of the Operational Research Society*, 50:167–176, 1999.

[24] V. Maniezzo and A. Colorni. The Ant System applied to the quadratic assignment problem. *IEEE Transactions on Data and Knowledge Engineering*, 11(5):769–778, 1999.

[25] B. Bullnheimer, R.F. Hartl, and C. Strauss. An improved ant system algorithm for the vehicle routing problem. *Annals of Operations Research*, 89:319–328, 1999.

[26] L.M. Gambardella, E.D. Taillard, and G. Agazzi. MACS-VRPTW: A multiple ant colony system for vehicle routing problems with time windows. In D. Corne, M. Dorigo, and F. Glover, editors, *New Ideas in Optimization*, pages 63–76. McGraw Hill, London, UK, 1999.

[27] C. Solnon. Solving permutation constraint satisfaction problems with artificial ants. In W. Horn, editor, *Proceedings of ECAI'2000, IOS Press, Amsterdam, The Netherlands*, pages 118–122, 2000.

[28] C. Solnon. Ants can solve constraint satisfaction problems. *IEEE Transactions on Evolutionary Computation*, 6(4):347–357, 2002.

[29] S. Fenet and C. Solnon. Searching for maximum cliques with ant colony optimization. In S. Cagnoni, J. Gottlieb, E. Hart, M. Middendorf, and G.R. Raidl, editors, *Applications of Evolutionary Computing, Proceedings of EvoWorkshops 2003: EvoCOP, EvoIASP, EvoSTim*, volume 2611 of LNCS, pages 236–245. Springer-Verlag, 2003.

[30] T. N. Bui and J. R. Rizzo. Finding maximum cliques with distributed ants. In K. Deb, R. Poli, W. Banzhaf, H.-G. Beyer, E.K. Burke, P.J. Darwen, D. Dasgupta, D. Floreano, J.A. Foster, M. Harman, O. Holland, P.L. Lanzi, L. Spector, A. Tettamanzi, D. Thierens, and A.M. Tyrell, editors, *Genetic and Evolutionary Computation - GECCO 2004*, volume 3102 of LNCS, pages 24–35. Springer-Verlag, 2004.

[31] J.I. van Hemert and T. Bäck. Measuring the searched space to guide efficiency: The principle and evidence on constraint satisfaction. In J.J. Merelo, A. Panagiotis, H.-G. Beyer, José-Luis Fernández-Villacañas, and Hans-Paul Schwefel, editors, *Proceedings of the 7th International Conference on Parallel Problem Solving from Nature*, volume 2439 of LNCS, pages 23–32, Berlin, 2002. Springer-Verlag.

[32] A. Colorni, M. Dorigo, and V. Maniezzo. An investigation of some properties of an ant algorithm. In R. Männer and B. Manderick, editors, *Parallel Problem Solving from Nature 2, PPSN-II, Brussels, Belgium*, pages 515–526. Elsevier, 1992.

[33] C. Solnon. Boosting ACO with a preprocessing step. In S. Cagnoni, J. Gottlieb, E. Hart, M. Middendorf, and G.R. Raidl, editors, *Applications of Evolutionary Computing, Proceedings of EvoWorkshops2002: EvoCOP, EvoIASP, EvoSTim*, volume 2279 of LNCS, pages 161–170. Springer-Verlag, 2002.

[34] M. Birattari, T. Stützle, L. Paquete, and K. Varrentrapp. A racing algorithm for configuring metaheuristics. In W. B. Langdon, E. Cantú-Paz,

K. Mathias, R. Roy, D. Davis, R. Poli, K. Balakrishnan, V. Honavar, G. Rudolph, J. Wegener, L. Bull, M. A. Potter, A. C. Schultz, J. F. Miller, E. Burke, and N. Jonoska, editors, *Proceedings of the Genetic and Evolutionary Computation Conference (GECCO-2002)*, pages 11–18. Morgan Kaufmann Publishers, San Mateo, CA, 2002.

[35] R. Burke, S. Gustafson, and G. Kendall. A survey and analysis of diversity measures in genetic programming. In W. B. Langdon, E. Cantú-Paz, K. Mathias, R. Roy, D. Davis, R. Poli, K. Balakrishnan, V. Honavar, G. Rudolph, J. Wegener, L. Bull, M. A. Potter, A. C. Schultz, J. F. Miller, E. Burke, and N. Jonoska, editors, *GECCO 2002: Proceedings of the Genetic and Evolutionary Computation Conference*, pages 716–723, New York, 9-13 July 2002. Morgan Kaufmann Publishers.

[36] J.I. van Hemert and C. Solnon. A study into ant colony optimization, evolutionary computation and constraint programming on binary constraint satisfaction problems. In S. Cagnoni, J. Gottlieb, E. Hart, M. Middendorf, and G.R. Raidl, editors, *Evolutionary Computation in Combinatorial Optimization (EvoCOP 2004)*, volume 3004 of LNCS, pages 114–123. Springer-Verlag, 2004.

[37] A. Sidaner, O. Bailleux, and J.-J. Chabrier. Measuring the spatial dispersion of evolutionist search processes: application to walksat. In P. Collet, C. Fonlupt, J.-K. Hao, E. Lutton, and M. Schoenauer, editors, *5th International Conference EA 2001*, volume 2310 of LNCS, pages 77–90. Springer-Verlag, 2001.

[38] R.W. Morrison and K.A. De Jong. Measurement of population diversity. In P. Collet, C. Fonlupt, J.-K. Hao, E. Lutton, and M. Schoenauer, editors, *5th International Conference EA 2001*, volume 2310 of LNCS, pages 31–41. Springer-Verlag, 2001.

[39] H.R. Lourenço, O. Martin, and T. Stützle. Iterated local search. In F. Glover and G. Kochenberger, editors, *Handbook of metaheuristics*, pages 321–353. Kluwer Academic Publishers, Dordrecht, The Netherlands, 2002.

[40] M. Brockington and J.C. Culberson. Camouflaging independent sets in quasi-random graphs. In D.S. Johnson and M.A. Trick, editors, *Cliques, Coloring, and Satisfiability*, volume 26 of *DIMACS Series in Discrete Mathematics and Theoretical Computer Science*, pages 75–88. American Mathematical Society, Boston, MA, USA, 1996.